

# Spotify Dataset Analysis

**Author: Mohan Krishna**

**This project analyzes a dataset of popular Spotify tracks using various statistical and machine learning techniques.**

**It includes data cleaning, exploratory data analysis, and a predictive model for track popularity.**

```
In [23]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import LinearSVC
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import cross_val_score
```

In [31]:

```
df = pd.read_csv('spotify_dataset.csv', encoding='ISO-8859-1')

df = df.drop(columns=['Unnamed: 0'])

df.rename(columns={'Track.Name': 'track_name', 'Artist.Name': 'artist_name', 'Beats Per Minute': 'beats_per_minute', 'Energy': 'Energy', 'Danceability': 'Danceability', 'Loudness..dB..': 'Loudness(dB)', 'Valence.': 'Valence', 'Length..': 'Length', 'Acousticness.': 'Acousticness', 'Speechiness.': 'Speechiness'})

df.head()
```

Out[31]:

	track_name	artist_name	Genre	beats_per_minute	Energy	Danceability	Loudness(dB)	Li
0	Señorita	Shawn Mendes	canadian pop	117	55	76	-6	
1	China	Anuel AA	reggaeton flow	105	81	79	-4	
2	boyfriend (with Social House)	Ariana Grande	dance pop	190	80	40	-4	
3	Beautiful People (feat. Khalid)	Ed Sheeran	pop	93	65	64	-8	
4	Goodbyes (Feat. Young Thug)	Post Malone	dfw rap	150	65	58	-4	

```
In [33]: # Check for missing values
print(df.isnull().sum())

# If there are any missing values, fill them with 0
df.fillna(0, inplace=True)

# Basic data statistics
print(df.describe())
```


```
track_name      0
artist_name     0
Genre           0
beats_per_minute 0
Energy          0
Danceability    0
Loudness(dB)    0
Liveness        0
Valence         0
Length          0
Acousticness    0
Speechiness     0
Popularity      0
dtype: int64
```

	beats_per_minute	Energy	Danceability	Loudness(dB)	Liveness \
count	50.000000	50.000000	50.00000	50.000000	50.000000
mean	120.060000	64.060000	71.38000	-5.660000	14.660000
std	30.898392	14.231913	11.92988	2.056448	11.118306
min	85.000000	32.000000	29.00000	-11.000000	5.000000
25%	96.000000	55.250000	67.00000	-6.750000	8.000000
50%	104.500000	66.500000	73.50000	-6.000000	11.000000
75%	137.500000	74.750000	79.75000	-4.000000	15.750000
max	190.000000	88.000000	90.00000	-2.000000	58.000000


  

	Valence	Length	Acousticness	Speechiness	Popularity
count	50.000000	50.000000	50.000000	50.000000	50.000000
mean	54.600000	200.960000	22.160000	12.480000	87.500000
std	22.336024	39.143879	18.995553	11.161596	4.491489
min	10.000000	115.000000	1.000000	3.000000	70.000000
25%	38.250000	176.750000	8.250000	5.000000	86.000000
50%	55.500000	198.000000	15.000000	7.000000	88.000000
75%	69.500000	217.500000	33.750000	15.000000	90.750000
max	95.000000	309.000000	75.000000	46.000000	95.000000


```
In [35]: # Correlation matrix to understand relationships between features
corr = df.corr(method='spearman')
plt.figure(figsize=(10, 10))
plt.title('Correlation Heatmap')
sns.heatmap(corr, annot=True, vmin=-1, vmax=1, cmap='GnBu_r', center=1)
plt.show()
```




```
In [37]: # Visualization of Genre distribution
genre_groupby = df.groupby('Genre')['track_name'].agg(len)
plt.figure(figsize=(15, 7))
xtick = ['dance pop', 'pop', 'latin', 'edm', 'canadian hip hop',
         'panamanian pop', 'electropop', 'reggaeton flow', 'canadian pop',
         'reggaeton', 'dfw rap', 'brostep', 'country rap', 'escape room',
         'trap music', 'big room', 'boy band', 'pop house', 'australian pop',
         'r&b en espanol', 'atl hip hop']
length = np.arange(len(xtick))
plt.bar(length, genre_groupby)
plt.xticks(length, xtick)
plt.xticks(rotation=90)
plt.ylabel('Count of Tracks', fontsize=20)
plt.show()
```




```
In [39]: # Number of songs by an artist
plt.figure(figsize=(10, 7))
df.groupby('artist_name')['track_name'].agg(len).sort_values(ascending=False).
plt.xlabel('Artist Name', fontsize=20)
plt.ylabel('Count of Songs', fontsize=20)
plt.title('Artist Name vs Count of Songs', fontsize=30)
plt.show()
```



```
In [41]: # Analyzing the relationship between Energy and Loudness
plt.figure(figsize=(10, 10))
sns.regplot(x='Energy', y='Loudness(dB)', data=df, color='black')
plt.title('Energy vs Loudness')
plt.show()
```



```
In [43]: # Pie chart of top artists by song count
labels = df.artist_name.value_counts().index
sizes = df.artist_name.value_counts().values
colors = ['red', 'yellowgreen', 'lightcoral', 'lightskyblue', 'cyan', 'green',
plt.figure(figsize=(10, 10))
plt.pie(sizes, labels=labels, colors=colors, autopct='%1.1f%%', startangle=140)
plt.axis('equal')
plt.title('Artist Song Distribution')
plt.show()
```



```
In [47]: # Train-test split for machine learning model
X = df.loc[:, ['Energy', 'Danceability', 'Length', 'Loudness(dB)', 'Acousticness']]
y = df.loc[:, 'Popularity'].values
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, random_state=42)
```

```
In [49]: # Linear regression model
regressor = LinearRegression()
regressor.fit(X_train, y_train)
```

```
Out[49]: LinearRegression
LinearRegression()
(https://scikit-learn.org/1.4/modules/generated/sklearn.linear_model.LinearRegression)
```




```
In [51]: # Predicted vs Actual values
y_pred = regressor.predict(X_test)
df_output = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
print(df_output)
```

	Actual	Predicted
0	89	87.638549
1	89	87.801842
2	88	87.504207
3	80	87.561699
4	85	88.098675
5	86	87.755355
6	88	87.495072
7	89	87.489517
8	88	87.430518
9	89	87.729801
10	89	87.531288
11	78	87.267113
12	91	88.293439
13	91	87.288000
14	91	88.107058

```
In [53]: # Error analysis
MSE = mean_squared_error(y_test, y_pred)
print('Mean Squared Error:', MSE)
```

Mean Squared Error: 13.02543638050829

```
In [55]: # Plotting Error Analysis
plt.figure(figsize=(10, 10))
plt.plot(y_pred, y_test, color='black', linestyle='dashed', marker='*', marker-
plt.title('Error Analysis')
plt.xlabel('Predicted Values')
plt.ylabel('Test Values')
plt.show()
```



```
In [57]: # Cross validation score
mse = cross_val_score(regressor, X_train, y_train, scoring='neg_mean_squared_e
mse_mean = np.mean(mse)
print('Cross-validated MSE:', abs(mse_mean))
```

Cross-validated MSE: 30.649516015364203

```
In [59]: # Gaussian Naive Bayes model
gnb = GaussianNB()
gnb.fit(X_train, y_train)
y_pred_gnb = gnb.predict(X_test)
```

```
In [61]: # Actual vs Predicted for Naive Bayes
df_gnb = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred_gnb})
print(df_gnb)
```


	Actual	Predicted
0	89	89
1	89	88
2	88	90
3	80	92
4	85	88
5	86	84
6	88	90
7	89	87
8	88	87
9	89	88
10	89	91
11	78	83
12	91	91
13	91	90
14	91	88

```
In [63]: # Support Vector Classifier model
LinSVC = LinearSVC(penalty='l2', loss='squared_hinge', dual=True)
LinSVC.fit(X_train, y_train)
y_pred_svc = LinSVC.predict(X_test)
```

```
In [65]: # Actual vs Predicted for SVC
df_svc = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred_svc})
print(df_svc)
```

	Actual	Predicted
0	89	95
1	89	84
2	88	84
3	80	84
4	85	84
5	86	84
6	88	84
7	89	84
8	88	84
9	89	95
10	89	84
11	78	95
12	91	70
13	91	84
14	91	84

```
In [67]: # Joint plot for prediction comparison  
sns.jointplot(x=y_test, y=y_pred, kind='reg', color='b')  
plt.show()
```



In [ ]: